

# Linear Dynamical Systems

SWC Neuroinformatics 2024

Dr. Aniruddh Galgali


Gatsby Unit, UCL

22 Feb 2024

# Last week

## Key assumptions


1. datapoints are independent & identically distributed (i.i.d)
2. given  $x$  (regressor), model dependence of  $y$  on  $x$  (“discriminative” model)



**Most datasets in neuroscience typically violate assumption (1)**

**Today** - a model for ‘more’ structured datasets (time-series)


# A day in the lab



# A day in the lab




overhead  
video  
camera




# A day in the lab



overhead  
video  
camera



Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$



Adjacent  $\mathbf{y}_t$  will be highly correlated (i.e not independent)

$\mathbf{y}_t$  is very high-dimensional with lots of redundant information

How can we leverage this to extract useful information?

# Outline


- Linear dynamical systems and state estimation
  - Applications in neuroscience
  - Model & problem definition
- Mathematical Preliminaries
  - A review of gaussian distributions
  - Bayes Rule
- Optimal state estimation - Kalman Filter
  - Prediction
  - Filtering
  - Smoothing

# Motivation

$\mathbf{y}_t$




overhead  
video  
camera




Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$


$\mathbf{y}_t =$




# Motivation



**Problem :** We don't observe  $X$  directly, but instead would like to infer/estimate it from  $Y$   
(State Estimation)



Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$



State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$

$$\mathbf{x}_t = \begin{bmatrix} x_t^1 \\ x_t^2 \end{bmatrix}$$

Position in  
2D space


**OR**

$$\mathbf{x}_t = \begin{bmatrix} x_t^1 \\ x_t^2 \\ \dot{x}_t^1 \\ \dot{x}_t^2 \end{bmatrix}$$

Position &  
Velocity in  
2D space




# The problem is ill-posed




Data:  $\mathbf{Y} = [y_1, y_2, \dots, y_T]$

State:  $\mathbf{X} = [x_1, x_2, \dots, x_T]$




# State Estimation requires probabilities




Data:  $\mathbf{Y} = [y_1, y_2, \dots, y_T]$

State:  $\mathbf{X} = [x_1, x_2, \dots, x_T]$



We would like to **infer** a probability distribution over possible states (locations) that the animal could be in.

# Application #2 – Brain Computer Interfaces



Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$


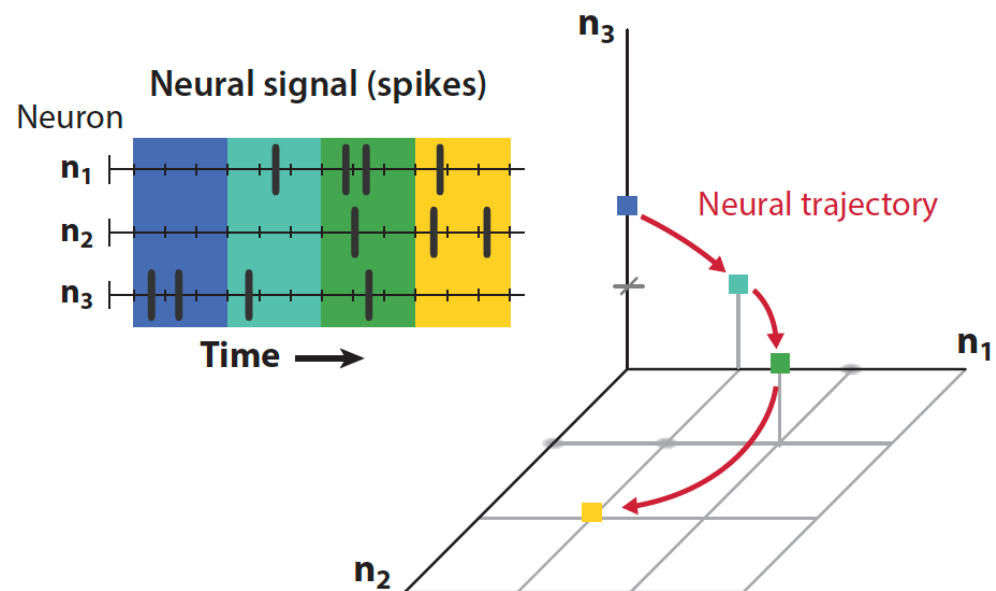
$$\mathbf{y}_t = \begin{bmatrix} \text{activity of neuron 1} \\ \text{activity of neuron 2} \\ \vdots \\ \text{activity of neuron 'N'} \end{bmatrix}$$

State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$


$$\mathbf{x}_t = \begin{bmatrix} \text{cursor vel. along } \rightarrow \\ \text{cursor vel. along } \uparrow \end{bmatrix}$$

**Goal:** Infer/Decode cursor locations (in real-time) based on spiking activity.

# Application #3 – Inferring population dynamics



Infer rules that govern evolution of  $X$  (dynamics)




Neural Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$

$$\mathbf{y}_t = \begin{bmatrix} \text{activity of neuron 1} \\ \text{activity of neuron 2} \\ \vdots \\ \text{activity of neuron 'N'} \end{bmatrix}$$

Population State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$

$$\mathbf{x}_t = \begin{bmatrix} x_t^1 \\ x_t^2 \end{bmatrix}$$

A 'low-d' representation that captures activity patterns across a population



# Specifying a model

pixels, neural activity, ...

Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$


$q$ -dimensional

location, cursor vel, neural pop. state

State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$

$p$ -dimensional

‘**Markov**’ assumption



“**Emission**” model – specifies how  $\mathbf{y}_t$  depends on  $\mathbf{x}_t$  (probabilistically)

$$p(\mathbf{y}_t | \mathbf{x}_t)$$

“**Dynamics**” model – specifies how  $\mathbf{x}_{t+1}$  depends on  $\mathbf{x}_t$

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t)$$

For a ‘linear’ dynamical system:

- Distributions are **gaussian**

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t+1}; \overset{\text{random variable}}{\mathbf{A}\mathbf{x}_t}, \overset{\text{mean}}{\mathbf{Q}})$$

$$p(\mathbf{y}_t | \mathbf{x}_t) = \mathcal{N}(\mathbf{y}_t; \mathbf{H}\mathbf{x}_t, \mathbf{R})$$

$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0) \quad \text{“Initial Condition”}$$

- Linear** in parameters

$\mathbf{A}, \mathbf{Q}, \mathbf{Q}_0$  –  $p \times p$  matrices

$\mathbf{H}$  –  $q \times p$  matrix

$\mathbf{R}$  –  $q \times q$  matrix

$\boldsymbol{\mu}_0$  –  $p \times 1$  vector

# A ‘generative’ model for time series

pixels, neural activity, ...

Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$

$q$ -dimensional


location, cursor vel, neural pop. state

State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$

$p$ -dimensional

‘Markov’ assumption

Inference  
 (“Bayes Rule”)



“Emission” model – specifies how  $\mathbf{y}_t$  depends on  $\mathbf{x}_t$  (probabilistically)

$$p(\mathbf{y}_t | \mathbf{x}_t)$$

“Dynamics” model – specifies how  $\mathbf{x}_{t+1}$  depends on  $\mathbf{x}_t$

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t)$$

For a ‘linear’ dynamical system:

- Distributions are **gaussian**

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t+1}; \overset{\text{random variable}}{\mathbf{A}\mathbf{x}_t}, \overset{\text{mean}}{\mathbf{Q}})$$

$$p(\mathbf{y}_t | \mathbf{x}_t) = \mathcal{N}(\mathbf{y}_t; \mathbf{H}\mathbf{x}_t, \mathbf{R})$$

$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0) \quad \text{“Initial Condition”}$$

- Linear in parameters

$\mathbf{A}, \mathbf{Q}, \mathbf{Q}_0$  –  $p \times p$  matrices

$\mathbf{H}$  –  $q \times p$  matrix

$\mathbf{R}$  –  $q \times q$  matrix

$\boldsymbol{\mu}_0$  –  $p \times 1$  vector

# Outline

- Linear dynamical systems and state estimation
  - Applications in neuroscience
  - Defining the model
- Mathematical Preliminaries (\*see whiteboard\*)
  - Bayes Rule
  - A review of gaussian distributions
- Optimal state estimation - Kalman Filter
  - Prediction
  - Filtering
  - Smoothing

# Alternative model definition

pixels, neural activity, ...


Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$

$q$ -dimensional ↙

location, cursor vel, neural pop. state

State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$

$p$ -dimensional ↙



**Follows from linearity and affine properties of gaussian distributions**

For a ‘linear’ dynamical system:

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t+1}; \overset{\text{random variable}}{\mathbf{A}} \mathbf{x}_t, \overset{\text{mean}}{\mathbf{Q}})$$

$$p(\mathbf{y}_t | \mathbf{x}_t) = \mathcal{N}(\mathbf{y}_t; \mathbf{H} \mathbf{x}_t, \mathbf{R})$$

$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0)$$



‘equivalent’

$$\begin{aligned} \mathbf{x}_{t+1} &= \mathbf{A} \mathbf{x}_t + \boldsymbol{\varepsilon}_t \\ \mathbf{y}_t &= \mathbf{H} \mathbf{x}_t + \boldsymbol{\eta}_t \end{aligned}$$

$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0)$$

$$p(\boldsymbol{\varepsilon}_t) = \mathcal{N}(\boldsymbol{\varepsilon}; \mathbf{0}, \mathbf{Q})$$

$$p(\boldsymbol{\eta}_t) = \mathcal{N}(\boldsymbol{\eta}; \mathbf{0}, \mathbf{R})$$



# Outline

- Linear dynamical systems and state estimation
  - Applications in neuroscience
  - Defining the model
- Mathematical Preliminaries (\*see whiteboard\*)
  - Bayes Rule
  - A review of gaussian distributions
- Optimal state estimation - Kalman Filter
  - Prediction
  - Filtering
  - Smoothing

# Optimal State Estimation


pixels, neural activity, ...

Data:  $\mathbf{Y} = [\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T]$   
*q-dimensional*

location, cursor vel, ...

State:  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$   
*p-dimensional*


Kalman Filter + Smoother



**Recursively** apply Bayes rule as new data come in

Decomposes into 3 fundamental steps

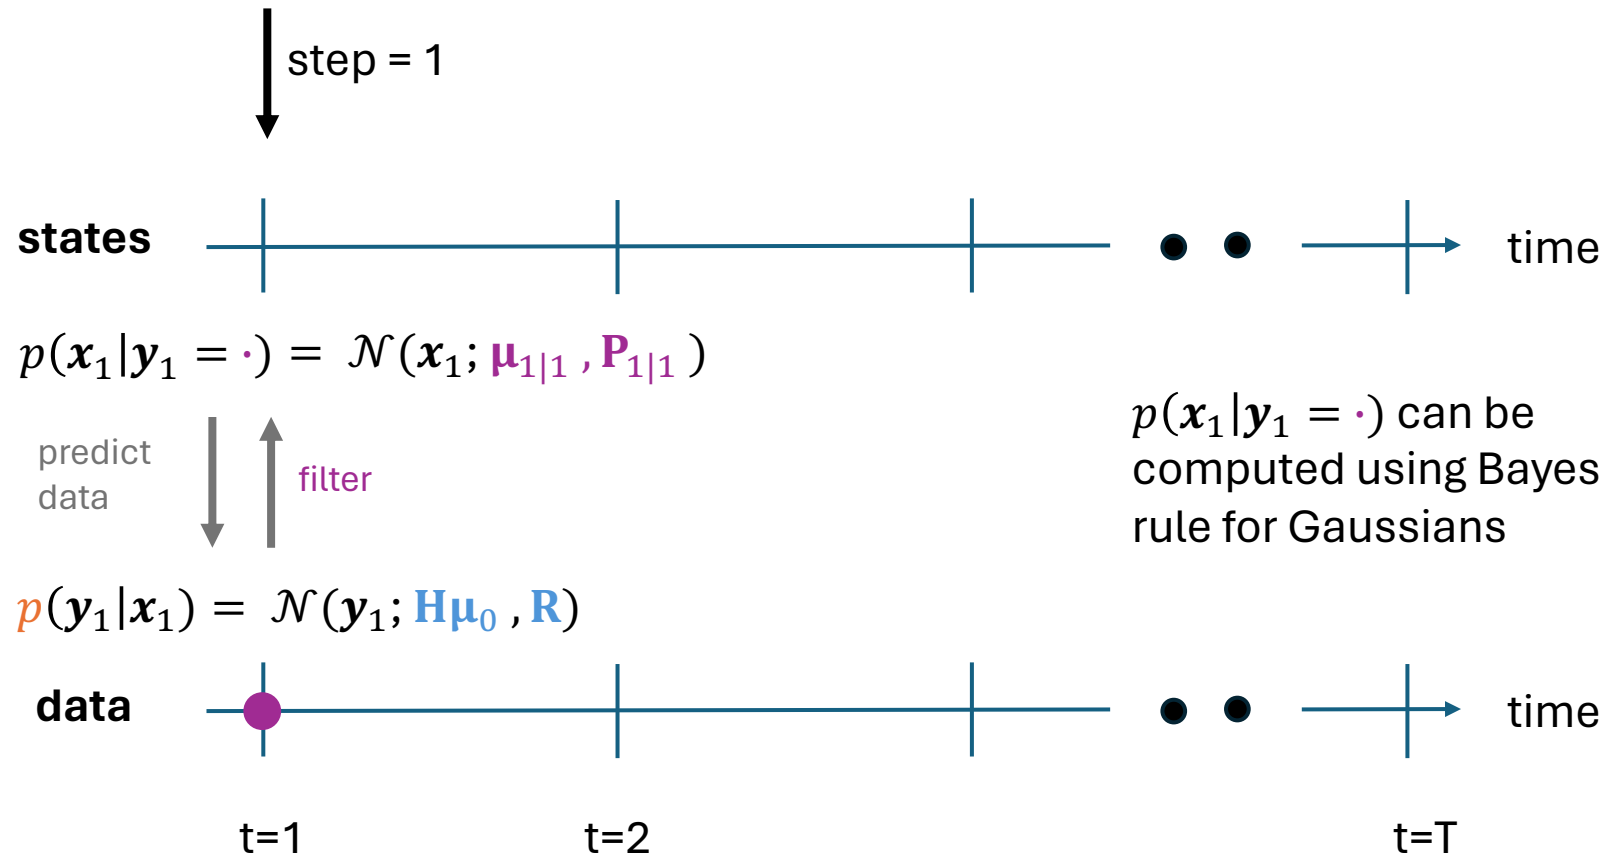
1. **Prediction** – Predict what the next best state is given the data I have observed so far
2. **Filtering** – Use disparity between model prediction and new measurement to update prediction
3. **Smoothing** – Use knowledge of entire sequence to further refine estimate of state.



Optimal in what sense ?

- **maximize probability** ('likelihood') of observing the data :  $P(\mathbf{y}_1, \dots, \mathbf{y}_T)$
- **minimize discrepancy** between LDS model prediction ( $\hat{\mathbf{y}}_t$ ) and true observation ( $\mathbf{y}_t$ ) but in a recursive least squares framework.

# Kalman prediction & filtering



## Model definition

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t+1}; \mathbf{A}\mathbf{x}_t, \mathbf{Q})$$

$$p(\mathbf{y}_t | \mathbf{x}_t) = \mathcal{N}(\mathbf{y}_t; \mathbf{H}\mathbf{x}_t, \mathbf{R})$$

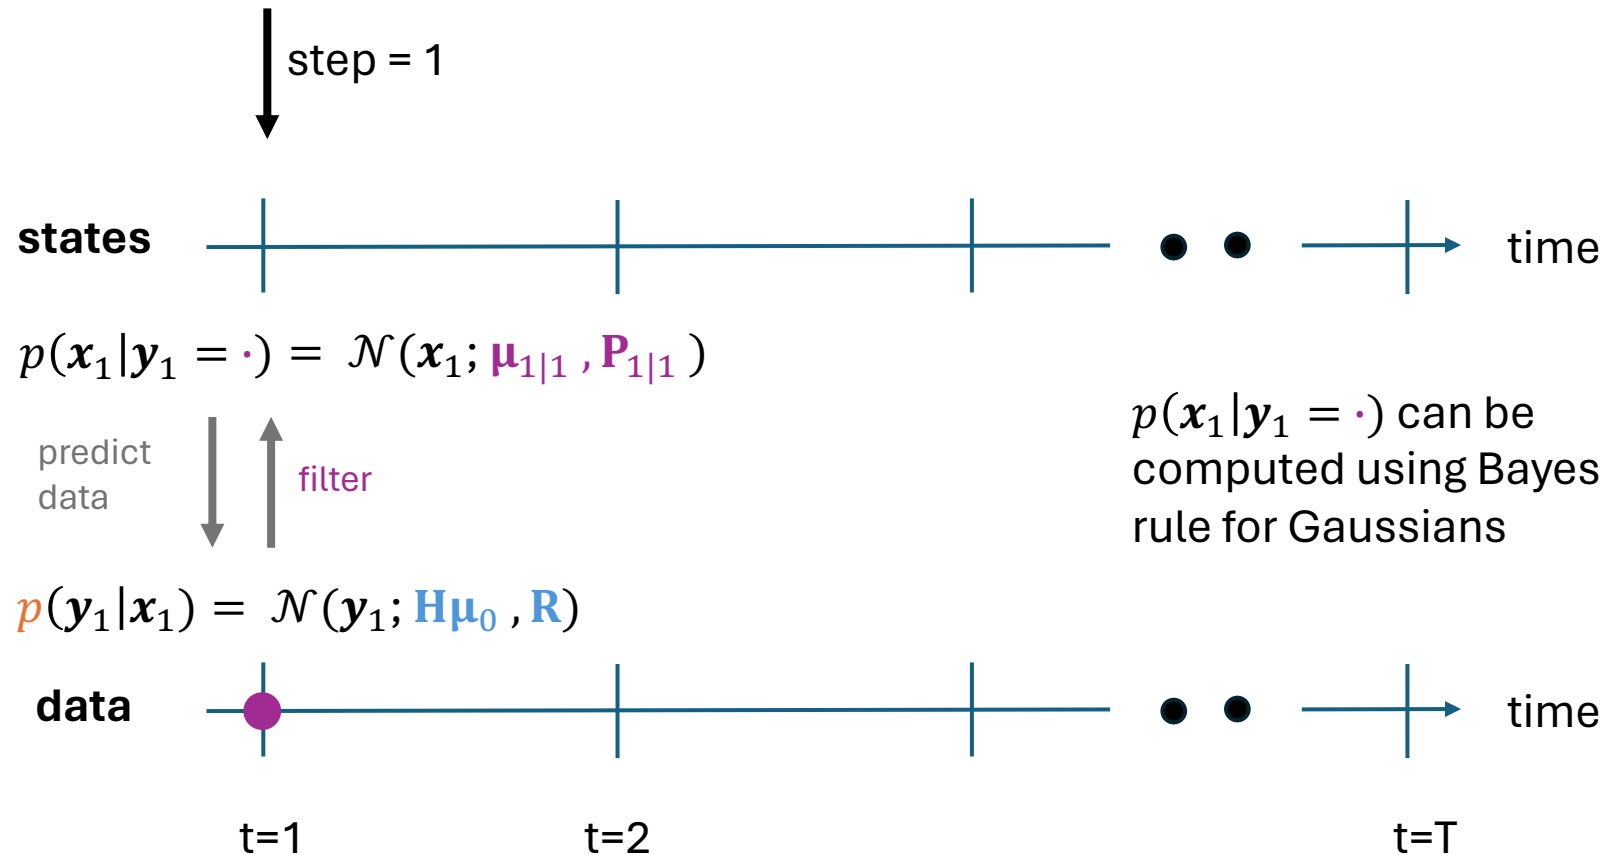
$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0)$$

## At current step:

$$\boldsymbol{\mu}_{1|1} = ?$$

$$\mathbf{P}_{1|1} = ?$$

# Kalman prediction & filtering



## Model definition

$$p(x_{t+1} | x_t) = \mathcal{N}(x_{t+1}; Ax_t, Q)$$

$$p(y_t | x_t) = \mathcal{N}(y_t; Hx_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$

## At current step:

'residual'  
(data - pred)


$$\mu_{1|1} = \mu_0 + K_1(y_1 - H\mu_0)$$

$$P_{1|1} = (I - K_1H) Q_0$$

$$K_1 = Q_0H'(HQ_0H' + R)^{-1}$$

Kalman 'gain'

# Kalman prediction & filtering



## Model definition

$$p(\mathbf{x}_{t+1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t+1}; \mathbf{A}\mathbf{x}_t, \mathbf{Q})$$

$$p(\mathbf{y}_t | \mathbf{x}_t) = \mathcal{N}(\mathbf{y}_t; \mathbf{H}\mathbf{x}_t, \mathbf{R})$$

$$p(\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_0, \mathbf{Q}_0)$$

## From previous step:

$$p(\mathbf{x}_1 | \mathbf{y}_1 = \cdot) = \mathcal{N}(\mathbf{x}_1; \boldsymbol{\mu}_{1|1}, \mathbf{P}_{1|1})$$


## At current step:

$$\boldsymbol{\mu}_{2|1} = ?$$

$$\mathbf{P}_{2|1} = ?$$

$$p(\mathbf{x}_2 | \mathbf{y}_1 = \cdot) = \int p(\mathbf{x}_2 | \mathbf{x}_1) p(\mathbf{x}_1 | \mathbf{y}_1 = \cdot) d\mathbf{x}_1 = \mathcal{N}(\mathbf{x}_2; \boldsymbol{\mu}_{2|1}, \mathbf{P}_{2|1})$$

# Kalman prediction & filtering



## Model definition

$$p(x_{t+1} | x_t) = \mathcal{N}(x_{t+1}; Ax_t, Q)$$

$$p(y_t | x_t) = \mathcal{N}(y_t; Hx_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$

## From previous step:

$$p(x_1 | y_1 = \cdot) = \mathcal{N}(x_1; \mu_{1|1}, P_{1|1})$$

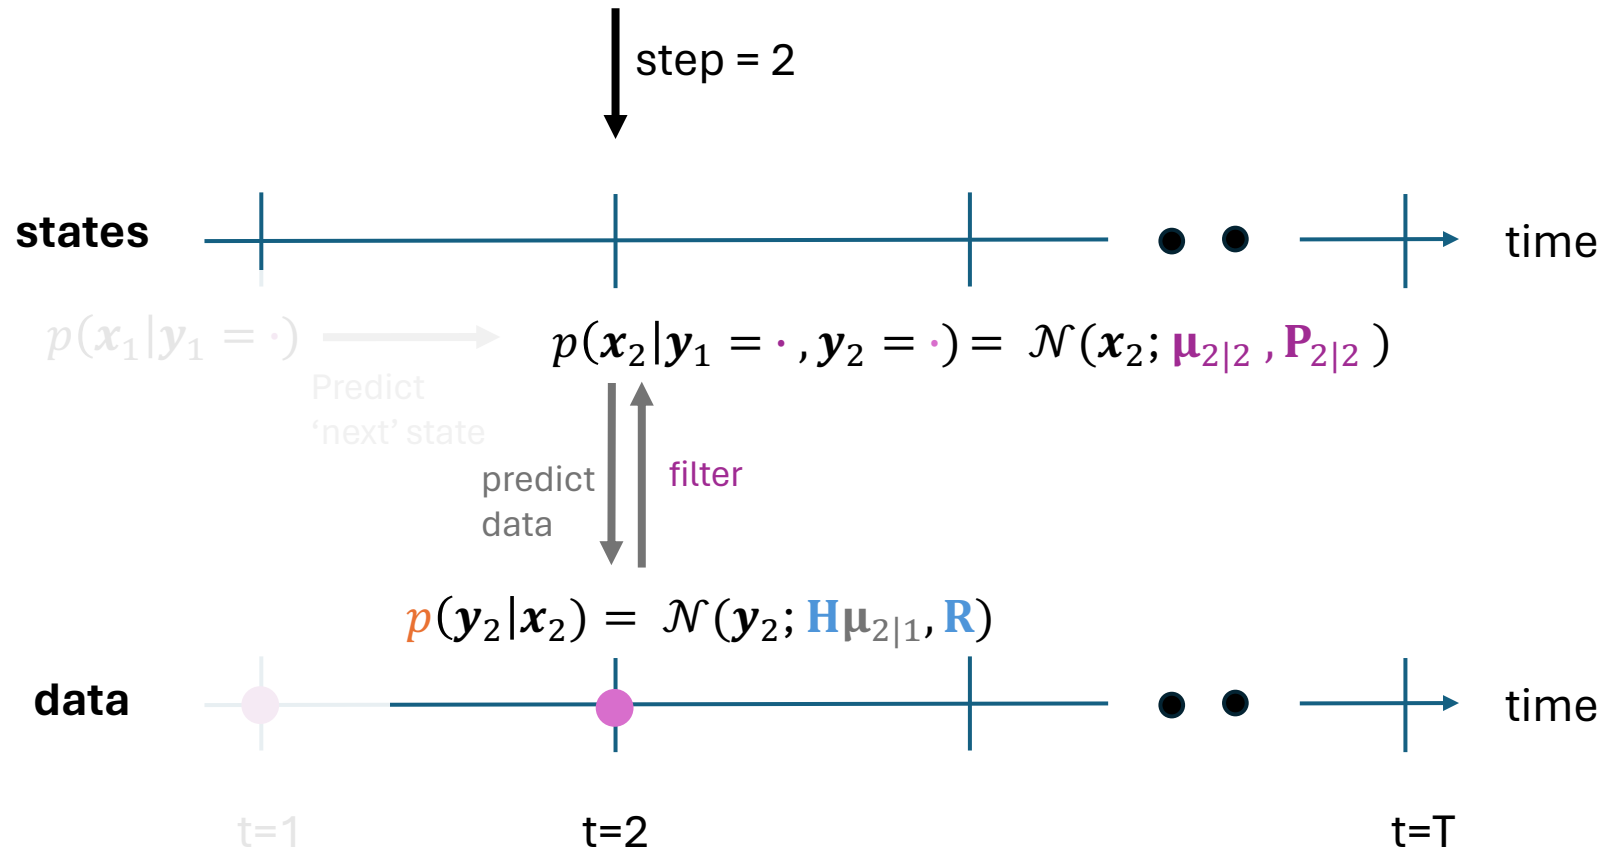
## At current step:

$$\mu_{2|1} = A\mu_{1|1}$$

$$P_{2|1} = AP_{1|1}A' + Q$$

$$p(x_2 | y_1 = \cdot) = \int p(x_2 | x_1) p(x_1 | y_1 = \cdot) dx_1 = \mathcal{N}(x_2; \mu_{2|1}, P_{2|1})$$

# Kalman prediction & filtering



## Model definition

$$p(x_{t+1}|x_t) = \mathcal{N}(x_{t+1}; Ax_t, Q)$$

$$p(y_t|x_t) = \mathcal{N}(y_t; Hx_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$

## From previous step:

$$p(x_1|y_1 = \cdot) = \mathcal{N}(x_1; \mu_{1|1}, P_{1|1})$$

## At current step:

$$\mu_{2|1} = A\mu_{1|1}$$

$$P_{2|1} = AP_{1|1}A' + Q$$

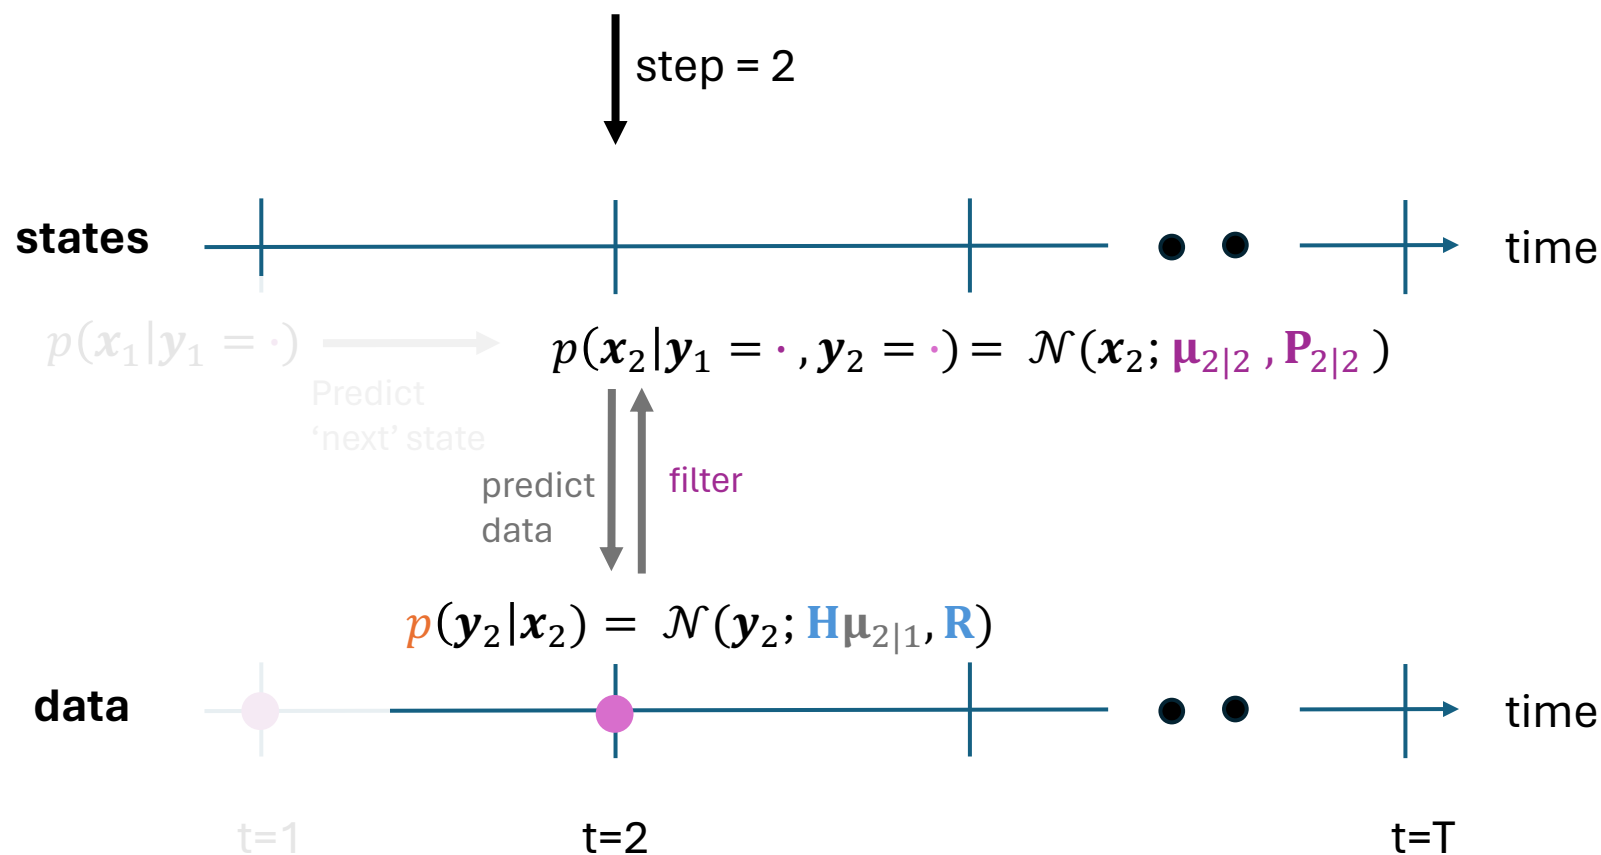
$$\mu_{2|2} = ?$$

$$P_{2|2} = ?$$

But 'predict data + filter' is analogous to what we did in Step 1!

'Initial' belief at step 2 is  $\mathcal{N}(x_2; \mu_{2|1}, P_{2|1})$  instead of  $\mathcal{N}(x_1; \mu_0, Q_0)$

# Kalman prediction & filtering



## Model definition

$$p(x_{t+1}|x_t) = \mathcal{N}(x_{t+1}; Ax_t, Q)$$

$$p(y_t|x_t) = \mathcal{N}(y_t; Hx_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$

## From previous step:

$$p(x_1|y_1 = \cdot) = \mathcal{N}(x_1; \mu_{1|1}, P_{1|1})$$

## At current step:

$$\mu_{2|1} = A\mu_{1|1}$$

$$P_{2|1} = AP_{1|1}A' + Q$$

$$\mu_{2|2} = \mu_{2|1} + K_2(y_2 - H\mu_{2|1})$$

$$P_{2|2} = (I - K_2H)P_{2|1}$$

$$K_2 = P_{2|1}H'(HP_{2|1}H' + R)^{-1}$$

But 'predict data + filter' is analogous to what we did in Step 1!

'Initial' belief at step 2 is  $\mathcal{N}(x_2; \mu_{2|1}, P_{2|1})$  instead of  $\mathcal{N}(x_1; \mu_0, Q_0)$



# Kalman prediction & filtering

Let,  $P(x_{t-1} | y_1, y_2, \dots, y_{t-1}) = \mathcal{N}(x_1; \mu_{t-1|t-1}, P_{t-1|t-1})$

1. Predict next state using 'dynamics' model:

$$\mu_{t|t-1} = A\mu_{t-1|t-1} \quad (\mu_{1|0} = \mu_0)$$

$$P_{t|t-1} = AP_{t-1|t-1}A' + Q \quad (P_{1|0} = Q_0)$$

2. Compute Kalman gain at 't'

$$K_t = P_{t|t-1}H'(HP_{t|t-1}H' + R)^{-1}$$

3. Update using observed data & 'emission' model

$$\mu_{t|t} = \mu_{t|t-1} + K_t(y_t - H\mu_{t|t-1})$$

$$P_{t|t} = (I - K_tH)P_{t|t-1}$$


**Online algorithm.** Really useful for real-time application (e.g BCIs, tracking)

## Model definition

$$p(x_{t+1} | x_t) = \mathcal{N}(x_{t+1}; Ax_t, Q)$$

$$p(y_t | x_t) = \mathcal{N}(y_t; Hx_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$



## Intuition

If model is very uncertain about measurements (i.e large  $R$  relative to  $Q$ ), then ignore data and use output of prediction as best state

# Kalman Smoothing

Let,  $P(x_{t+1} | y_1, y_2, \dots, y_T) = \mathcal{N}(x_{t+1}; \mu_{t+1|T}, P_{t+1|T})$

1. Start at  $t = T$  and **work backwards**

2. Compute reverse Kalman gain at ' $t$ '

$$J_t = P_{t|t} A' (P_{t|t-1})^{-1}$$

3. Update

$$\mu_{t|T} = \mu_{t|t} + J_t (\mu_{t+1|T} - \mu_{t+1|t})$$

$$P_{t|T} = P_{t|t} + J_t (P_{t+1|T} - P_{t+1|t}) J_t'$$


**Offline algorithm.** Can be computed only once entire data has been using. Useful for refinement.

**Model definition**

$$p(x_{t+1} | x_t) = \mathcal{N}(x_{t+1}; A x_t, Q)$$

$$p(y_t | x_t) = \mathcal{N}(y_t; H x_t, R)$$

$$p(x_1) = \mathcal{N}(x_1; \mu_0, Q_0)$$



Just need the quantities that were computed in the forward direction (prediction + filtering)

# Summary

- Time-series data are ubiquitous in neuroscience.
- An important problem – state estimation
  - Many applications – behavior tracking, neural/BCI decoding, inferring internal states, and many more..
  - Can be formulated as a recursive Bayesian estimation problem
- Bayesian state estimation is a **general** algorithm
  - For fixed model parameters, estimate distribution over states given data.
  - We looked at a specific case – linear, gaussian model – gives us an analytical solution (Kalman filtering + smoothing)
  - Dynamics/Emission models can be more general (think neural networks) – this requires sophisticated algorithms that approximate Bayes Rule.

# Further Reading

## Concepts

- Pattern Recognition & Machine Learning (Chris Bishop)
  - Chapter 2 – for gaussian distributions
  - Chapter 13 – for linear dynamical systems//hidden markov models
- A nice visualization of Gaussians and associated concepts
  - <https://distill.pub/2019/visual-exploration-gaussian-processes/>

## Applications:

- <https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python> (code notebook)
- Macke et al, NeurIPS 2013, Empirical models of spiking in neural populations
- Gilja et al, Nat. Neuro 2012, A high performance neural prosthesis enabled by control algorithm design.